Trees

Reference: ISL Ch. 8

Decision Trees: Using tree-logic to make predictions.

CART: Classification and Regression Trees.

Random Forests: averaging over many possible trees.

Simple examples: NBC, prostate cancer, motorcycle crashes.

Larger example on house prices in california.

What is a Decision Tree?

>70% Rain ──────────────────────→ Umbrella

Wake Up

>70% Rain

<70% Rain

>30% Rain

Cloudy ──→ Umbrella

Sunny ──→ No Umbrella

<30% Rain ──────────────→ No Umbrella

Tree-logic uses a series of steps to come to a conclusion.
The trick is to have mini-decisions combine for good choices.
Each decision is a node, and the final prediction is a leaf node

Decision Trees are a Regression Model

You have inputs **x** (forecast, current conditions)
and an output of interest y (need for an umbrella).

Based on previous data, the goal is to specify branches of
choices that lead to good predictions in new scenarios.
In other words, you want to estimate a Tree Model.

Instead of linear coefficients, we need to find 'decision nodes':
  split-rules defined via thresholds on some dimension of **x**.

Nodes have a parent-child structure: every node except the root
has a parent, and every node except the leaves has two children.

Decision trees are like a game of mousetrap

You drop your **x** covariates in at the top, and each decision node bounces you either left or right. Finally, you end up in a leaf node which contains the data subset defined by these decisions (splits).

$$x_i = 0$$

$$\swarrow \qquad \searrow$$

$$x_j = 2 \qquad \{\mathbf{x} : x_i > 0\}$$

$$\swarrow \qquad \searrow$$

$$\{\mathbf{x} : x_i \leq 0, x_j \leq 2\} \qquad \{\mathbf{x} : x_i \leq 0, x_j > 2\}$$

The prediction rule at each leaf (a class probability or predicted $\hat{y}$) is the average of the sample y values that ended up in that leaf.

Estimation of Decision Trees

As usual, we'll maximize data likelihood (minimize deviance).
But what are the observation probabilities in a tree model?

Two types of likelihood: classification and regression trees.

A given covariate **x** dictates your path through
tree nodes, leading to a leaf node at the end.

Classification trees have class probabilities at the leaves.
Probability I'll be in heavy rain is 0.9 (so take an umbrella).

Regression trees have a mean response at the leaves.
The expected amount of rain is 2in (so take an umbrella).

Tree deviance is the same as in linear models

Regression Deviance: $\sum_{i=1}^{n}(y_i - \hat{y}_i)^2$

Classification Deviance: $-\sum_{i=1}^{n}\log(\hat{p}_{y_i})$

It is also common to use Gini Deviance, $-\sum_{i=1}^{n}\hat{p}_{y_i}(1 - \hat{p}_{y_i})$

Instead of being based on $\mathbf{x}'\boldsymbol{\beta}$, predicted $\hat{p}$ and $\hat{y}$ are functions of $\mathbf{x}$ passed through the decision nodes.

We need a way to estimate the sequence of decisions.

▶ How many are they? What is the order?

There is a *huge* set of possible tree configurations.

Given a parent set of data $\{\mathbf{x}_i, y_i\}_{i=1}^n$, the optimal split is that location x_{ij} on some dimension j on some observation i, so that the child sets

$$\text{left: } \{\mathbf{x}_k, y_k : x_{kj} \leq x_{ij}\} \text{ and right: } \{\mathbf{x}_k, y_k : x_{kj} > x_{ij}\}$$

are as homogeneous in response y as possible.

For example, we will minimize the sum of squared errors

$$\sum_{k \in \text{left}} (y_k - \bar{y}_{\text{left}})^2 + \sum_{k \in \text{right}} (y_k - \bar{y}_{\text{right}})^2$$

for *regression trees*, or gini impurity for *classification trees* (e.g., the sum across children 'c' of $n_c \bar{y}_c (1 - \bar{y}_c)$ if $y \in \{0, 1\}$).

We estimate decision trees by being recursive and greedy

CART grows the tree through a sequence of splits:

- ▶ Given any set (node) of data, you can find the optimal split (the error minimizing split) and divide into two child sets.
- ▶ We then look at each child set, and again find the optimal split to divide it into two homogeneous subsets.
- ▶ The children become parents, and we look again for the optimal split on their new children (the grandchildren!).

You stop splitting and growing when the size of the leaf nodes hits some minimum threshold (e.g., say no less than 10 obsv per leaf). Often there are also minimum deviance improvement thresholds.

Use the tree library for CART in R

The syntax is essentially the same as for `glm`:

`mytree = tree(y ~x1 + x2 + x3 ..., data=mydata)`

There are only a few other possible arguments,
all of which dictate possible types of new children

- `mincut` is the minimum size for a new child.
- `mindev` is the minimum (proportion) deviance improvement
  for proceeding with a new split.

These are important: you may want to make them smaller
than their defaults: `mincut=5`, `mindev=0.01`.

As usual, you can `print`, `summarize`, and `plot` the tree.

Recall our NBC data: viewer demographic % by show.

Consider a classification tree to predict genre from demographics.

```
> genretree
node), split, n, deviance, yval, (yprob)
      * denotes terminal node

1) root 40 75.800 Drama/Adventure ( 0.475 0.425 0.100 )
 2) CABLE.W.O.PAY < 28.6651 22 33.420 Drama/Adventure (0.73 0.09 0.18)
  4) VCR.OWNER < 83.749 5  6.730 Situation Comedy (0.00 0.40 0.60) *
  5) VCR.OWNER > 83.749 17  7.606 Drama/Adventure (0.941 0.000 0.059)
   10) TERRITORY.EAST.CENTRAL < 16.4555 16  0.000 Drama/Adventure(1 0 0) *
   11) TERRITORY.EAST.CENTRAL > 16.4555 1  0.000 Situation Comedy(0 0 1) *
  3) CABLE.W.O.PAY > 28.6651 18 16.220 Reality ( 0.16667 0.83333 0 )
    6) BLACK < 17.2017 15  0.000 Reality ( 0 1 0 ) *
    7) BLACK > 17.2017 3  0.000 Drama/Adventure ( 1 0 0 ) *
```

Output from `tree` shows a series of decision nodes and the
proportion in each genre at these nodes, down to the leaves.

Trees are easiest to understand in a dendrogram

WIRED.CABLE.W.O.PAY < 28.6651

VCR.OWNER < 83.749

BLACK < 17.2017

Situation Comedy

TERRITORY.EAST.CENTRAL < 16.4555

Reality

Drama/Adventure

Drama/Adventure   Situation Comedy
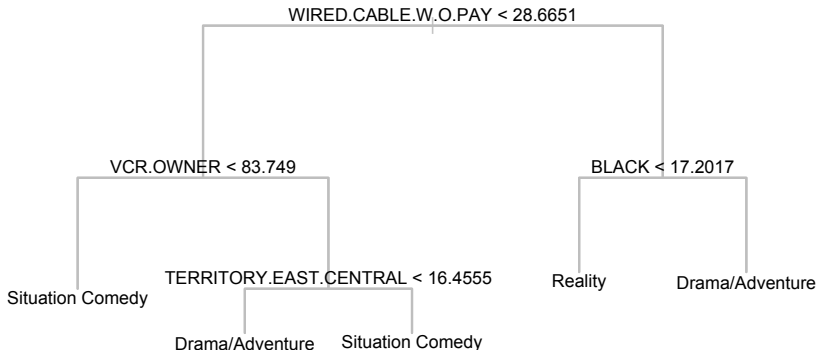
Shows the sequence of internal splits, ending in leaf-node decisions.
Here, the decision is the genre of highest proportion in each leaf.
To get the dendrogram, do `plot(mytree)` then `text(mytree)`.

Consider predicting engagement from ratings and genre.

Split on genre by turning it into a group of numeric variables:

```
x <- model.matrix(PE ~ Genre + GRP, data=nbc)[,-1]
names(x) <- c("reality","comedy","GRP")
```

We have a reference factor level (drama)

A regression tree:

```
nbctree <- tree(PE ~ ., data=x, mincut=1)
```

Instead of genre, leaf predictions are expected engagement.

mincut=1 allows for leaves containing a single show,
with expected engagement that single show's PE.

An NBC Show Engagement Tree
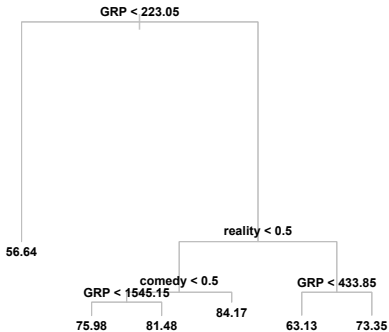
Green is comedy, blue is drama, red is reality

Nonlinear: PE increases with GRP, but in jumps

Follow how the tree translates into changing $\mathbb{E}[\text{PE}]$

Trees provide Automatic Interaction Detection

For example, different genres are more/less dependent on GRP.

AID was an original motivation for building decision trees.
Older algorithms have it in their name: CHAID, US-AID, ...

This is pretty powerful technology: nonlinearity and
interaction without having to specify it in advance.
Moreover, nonconstant variance is no problem.

Methods with these characteristics are called nonparametric.
No assumed parametric model (eg, $y = x\beta + \varepsilon, \ \varepsilon \sim \mathrm{N}(0, \sigma^2)$).

Pruning your tree for cross validation

Biggest challenge with such flexible models is avoiding overfit.
For CART, the usual solution is to rely on cross validation.

The basic constraints (`mincut, mindev`) lead to a full tree fit.
Prune this tree by removing split rules from the bottom up:

At each step, remove the split that contributes least to
deviance reduction, thus reversing CART's growth process.

Pruning yields candidate trees, and we use CV to choose.
Each prune step produces a candidate tree model, and we
can compare their out-of-sample prediction performance.

Example: Prostate Cancer Prognosis

After tumor detection, there are many treatment options.

▶ Various chemo + radiation, surgical removal.

Biopsy information is available to help in deciding treatment

▶ Gleason Score: microscopic pattern classes.

▶ Prostate Specific Antigen: protein production.

▶ Capsular Penetration: reach of cancer into gland lining.

▶ Benign Prostatic Hyperplasia Amount: size of prostate.

Another influential variable is the patient's age.

The goal is to predict tumor log-volume (size, spread).

Full tree fit to 97 prostate cancer patients

Leaf node labels are expected tumor `log(volume)`.
Do we need all the splits? Is the tree just fitting noise?

Cross-Validated Tree Pruning

cv.tree does cross-validation across pruning levels.
```
cvpst <- cv.tree(pstree, K=90) # K is nfolds
```


The output can be plotted, and it holds out-of-sample deviance for each tree size (the number of leaf nodes).

Pruning the Prostate Cancer Tree

Out-of-sample deviance can be used to choose tree size.

```
> cvpst$size
 [1] 12 11  8  7  6  5  4  3  2  1
> cvpst$dev
 [1] 72  75  77  76  76  77  77  70  97 160
```

Since size 3 has lowest CV deviance, it is 'best'.

To fit this tree, use the prune.tree function:

```
pstcut <- prune.tree(pstree, best=3)
```

pstcut is then itself a new tree object.

The Pruned Treatment Tree

CV chooses PSA and penetration as deciding variables.
Note the interaction: penetration effect depends on PSA.

Prostate Cancer Prognosis Tree

With only 2 relevant inputs, we can plot the data and tree fit.
Points proportional to tumor size, leaf partitions are in blue.

Trees are awesome

They automatically learn non-linear response functions and will discover interactions between variables.

Example: Motorcycle Crash Test Dummy Data
x is time from impact, y is acceleration on the helmet.

Unfortunately, it is tough to avoid overfit with CART:

Deep tree structure is so unstable that optimal depth is not easily chosen via cross validation, and there's no theory to fall back on.

Instead, we can average over a bootstrapped sample of trees:

- ▶ repeatedly re-sample the data, with-replacement,
  to get a 'jittered' dataset of n observations.
- ▶ for each resample, fit a CART tree.
- ▶ when you want to predict y for some $\mathbf{x}$,
  take the average prediction from this forest of trees.

Real structure that persists across datasets shows up in the average. Noisy useless signals will average out to have no effect.

This is a Random Forest

Random Forests

- Sample B subsets of the data + variables:
  e.g., observations $1, 5, 20, ...$ and inputs $2, 10, 17, ...$

- Fit a tree to each subset, to get B fitted trees is $\mathcal{T}_b$.

- Average prediction across trees:
  - for regression average $\mathbb{E}[y|\mathbf{x}] = \frac{1}{B} \sum_{b=1}^{B} \mathcal{T}_b(\mathbf{x})$.
  - for classification let $\{\mathcal{T}_b(\mathbf{x})\}_{b=1}^{B}$ vote on $\hat{y}$.

The observation resample is usually *with-replacement*, so that this is taking the *average of bootstrapped trees* (i.e., 'bagging')

Understanding Random Forests

Recall how CART is used in practice.

- ▶ Split to lower deviance until leaves hit minimum size.
- ▶ Create a set of candidate trees by pruning back from this.
- ▶ Choose the best among those trees by cross validation.

Random Forests avoid the need for CV.

Each tree 'b' is not overly complicated because
you only work with a limited set of variables.

Your predictions are not 'optimized to noise' because
they are averages of trees fit to many different subsets.

RFs are a great go-to model for nonparametric prediction.

Model Averaging

This technique of 'Model Averaging' is central to
many advanced nonparametric learning algorithms.

ensemble learning, mixture of experts, Bayesian averages, ...

It works best with flexible but simple models

Recall lasso as a stabilized version of stepwise regression
(if you jitter the data your estimates stay pretty constant).

Model averaging is a way to take arbitrary *unstable* methods, and
make them stable. This makes training easier.

Probability of rain on a new day is the average P(rain)
across some trees that split on forecast, others on sky.

We don't get tied to one way of deciding about umbrellas.

Random Forests in R

R has the randomForest package,
which works essentially the same as tree

```
spamrf <- randomForest(spam ~ ., data=xemail)
```

For big datasets, use $x = x$, $y = y$ like in gamlr.

Unfortunately, you lose the interpretability of a single tree.
However, if you set importance=TRUE, Random Forest will
evaluate each $\mathcal{T}_b$'s performance on the *left-out sample* (recall each
tree is fit on a sub-sample). This yields nice OOS stats.

They can be slow (due to many tree fits) but
they can also be fit in parallel or on distributed data...

Fitting CART to sampled-with-replacement data is equivalent to randomly weighting your observations in the deviance calculations.

(size proportional to weight in this picture).

Random Trees for the Motorcycle Data

If you fit to random subsets of the data,

you get a slightly different tree each time.

Model Averaging with Random Forests

Averaging many trees yields a single response surface.

Still looks like a bit of overfit to me, which remains a danger.

A larger example: California Housing Data

Median home values in census tracts, along with

- ▶ Latitude and Longitude of tract centers.
- ▶ Population totals and median income.
- ▶ Average room/bedroom numbers, home age.

The goal is to predict log(MedVal) for census tracts.

Difficult regression: Covariate effects change with location.
How they change is probably not linear.
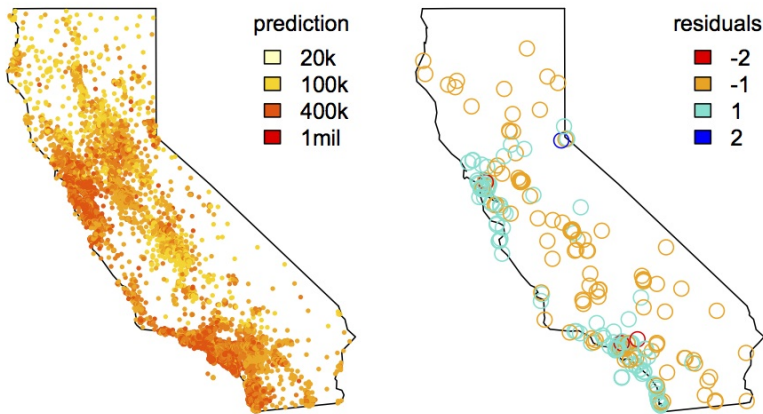
CART Dendrogram for CA housing

Income is dominant, with location important for low income.
Cross Validation favors the most complicated tree: 12 leaves.

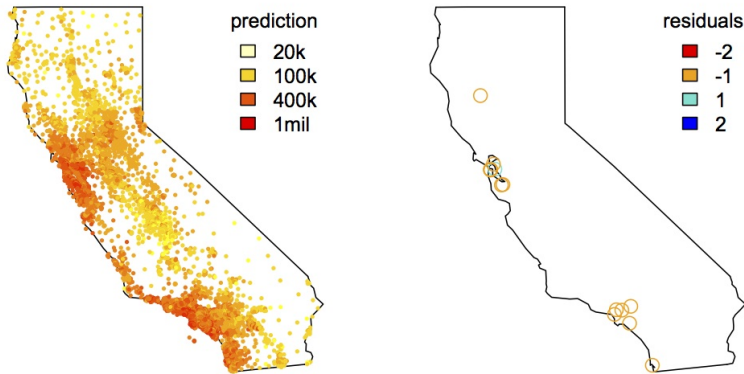**LASSO** **fit for CA housing data**

Looks like over-estimates in the Bay, under-estimates in OC.

CART fit for CA housing data

Under-estimating the coast, over-estimating the central valley?

randomForest fit for CA housing data

No big residuals! (although still missing the LA and SF effects)
Overfit? From out-of-sample prediction it appears not.

CA housing: out-of-sample prediction

Trees outperform LASSO: gain from nonlinear interaction.
RF is better still than CART: benefits of model averaging.

Although you don't have a nice single tree to interpret,
randomForest provides OOS variable importance plots.

You need to run randomForest with importance=TRUE. Otherwise
it doesn't store the necessary information.

RF variable importance

The x-axis here is the % amount that removing
  splits on that variable would increase the MSE.
For classification it plots increase in % misclassified.

Roundup on Tree-based learning

We've seen two techniques for building tree models.

- CART: recursive partitions, pruned back by CV.
- randomForest: average many simple CART trees.

There are many other tree-based algorithms.

- Boosted Trees: repeatedly fit simple trees to residuals.
  Fast, but it is tough to avoid over-fit (requires full CV).
- Bayes Additive Regression Trees: mix many simple trees.
  Robust prediction, but suffers with non-constant variance.
- Dynamic Trees: grow sequential 'particle' trees
  Good online, but fit depends on data ordering

Trees are poor in high dimension, but fitting them to low
dimension factors (principle components) is a good option.

Roundup on Nonlinear Regression and Classification

Many other *nonparametric learning* algorithms

- ▶ Neural Networks (and deep learning):
  many recursive logistic regressions.

- ▶ Support Vector Machines:
  Project to HD, then classify.

- ▶ Gaussian Processes, splines, wavelets, etc:
  Use sums of curvy functions in regression.

Some of these are great, but all take a ton of tuning.

Nothing's better out-of-the-box in low dimension than trees. But: when the (simpler) linear model fits, it will do better.
This is most often the case in very high dimension.